

```
import pandas as pd
import numpy as np
df=pd.read_csv('/kaggle/input/phishing-website-detector/phishing.csv')
df
```

	Index	UsingIP	LongURL	ShortURL	Symbol@	Redirecting//	\
0	0	1	1	1	1	1	
1	1	1	0	1	1	1	
2	2	1	0	1	1	1	
3	3	1	0	-1	1	1	
4	4	-1	0	-1	1	-1	
...
11049	11049	1	-1	1	-1	1	
11050	11050	-1	1	1	-1	-1	
11051	11051	1	-1	1	1	1	
11052	11052	-1	-1	1	1	1	
11053	11053	-1	-1	1	1	1	

	PrefixSuffix- UsingPopupWindow	SubDomains	HTTPS	DomainRegLen	...
0	-1	0	1	-1	...
1	-1	-1	-1	-1	...
1	-1	-1	-1	1	...
2	-1	-1	-1	-1	...
3	-1	1	1	-1	...
-1	-1	1	1	-1	...
4	-1	1	1	-1	...
1
...
11049	1	1	1	-1	...
-1	-1	1	-1	-1	...
11050	-1	1	-1	-1	...
-1	-1	1	-1	-1	...
11051	-1	-1	-1	1	...
-1	-1	-1	-1	1	...
11052	-1	-1	-1	1	...
-1	-1	-1	-1	1	...
11053	-1	-1	-1	1	...
1					

	IframeRedirection PageRank	AgeofDomain	DNSRecording	WebsiteTraffic
0	1	-1	-1	0
-1	1	1	-1	1
1	1	-1	-1	1
-1				
2				

-1				
3	1	-1	-1	0
-1				
4	1	1	1	1
-1				
...
...				
11049	-1	1	1	-1
-1				
11050	1	1	1	1
1				
11051	1	1	1	1
-1				
11052	1	1	1	1
-1				
11053	1	-1	1	-1
-1				

	GoogleIndex	LinksPointingToPage	StatsReport	class
0	1	1	1	-1
1	1	0	-1	-1
2	1	-1	1	-1
3	1	1	1	1
4	1	-1	-1	1
...
11049	1	1	1	1
11050	1	-1	1	-1
11051	1	0	1	-1
11052	1	1	1	-1
11053	-1	1	-1	-1

[11054 rows x 32 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 11054 entries, 0 to 11053

Data columns (total 32 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Index	11054 non-null	int64
1	UsingIP	11054 non-null	int64
2	LongURL	11054 non-null	int64
3	ShortURL	11054 non-null	int64
4	Symbol@	11054 non-null	int64
5	Redirecting//	11054 non-null	int64
6	PrefixSuffix-	11054 non-null	int64
7	SubDomains	11054 non-null	int64
8	HTTPS	11054 non-null	int64
9	DomainRegLen	11054 non-null	int64

10	Favicon	11054	non-null	int64
11	NonStdPort	11054	non-null	int64
12	HTTPSDomainURL	11054	non-null	int64
13	RequestURL	11054	non-null	int64
14	AnchorURL	11054	non-null	int64
15	LinksInScriptTags	11054	non-null	int64
16	ServerFormHandler	11054	non-null	int64
17	InfoEmail	11054	non-null	int64
18	AbnormalURL	11054	non-null	int64
19	WebsiteForwarding	11054	non-null	int64
20	StatusBarCust	11054	non-null	int64
21	DisableRightClick	11054	non-null	int64
22	UsingPopupWindow	11054	non-null	int64
23	IframeRedirection	11054	non-null	int64
24	AgeofDomain	11054	non-null	int64
25	DNSRecording	11054	non-null	int64
26	WebsiteTraffic	11054	non-null	int64
27	PageRank	11054	non-null	int64
28	GoogleIndex	11054	non-null	int64
29	LinksPointingToPage	11054	non-null	int64
30	StatsReport	11054	non-null	int64
31	class	11054	non-null	int64

dtypes: int64(32)

memory usage: 2.7 MB

df.describe()

	Index	UsingIP	LongURL	ShortURL	
Symbol@ \					
count	11054.000000	11054.000000	11054.000000	11054.000000	
11054.000000					
mean	5526.500000	0.313914	-0.633345	0.738737	
0.700561					
std	3191.159272	0.949495	0.765973	0.674024	
0.713625					
min	0.000000	-1.000000	-1.000000	-1.000000	-
1.000000					
25%	2763.250000	-1.000000	-1.000000	1.000000	
1.000000					
50%	5526.500000	1.000000	-1.000000	1.000000	
1.000000					
75%	8289.750000	1.000000	-1.000000	1.000000	
1.000000					
max	11053.000000	1.000000	1.000000	1.000000	
1.000000					
Redirecting//		PrefixSuffix-	SubDomains	HTTPS	
DomainRegLen \					
count	11054.000000	11054.000000	11054.000000	11054.000000	
11054.000000					

mean	0.741632	-0.734938	0.064049	0.251040	-
0.336711					
std	0.670837	0.678165	0.817492	0.911856	
0.941651					
min	-1.000000	-1.000000	-1.000000	-1.000000	-
1.000000					
25%	1.000000	-1.000000	-1.000000	-1.000000	-
1.000000					
50%	1.000000	-1.000000	0.000000	1.000000	-
1.000000					
75%	1.000000	-1.000000	1.000000	1.000000	
1.000000					
max	1.000000	1.000000	1.000000	1.000000	
1.000000					

	...	UsingPopupWindow	IframeRedirection	AgeofDomain	
DNSRecording \					
count	...	11054.000000	11054.000000	11054.000000	
11054.000000					
mean	...	0.613353	0.816899	0.061335	
0.377239					
std	...	0.789845	0.576807	0.998162	
0.926158					
min	...	-1.000000	-1.000000	-1.000000	-
1.000000					
25%	...	1.000000	1.000000	-1.000000	-
1.000000					
50%	...	1.000000	1.000000	1.000000	
1.000000					
75%	...	1.000000	1.000000	1.000000	
1.000000					
max	...	1.000000	1.000000	1.000000	
1.000000					

	WebsiteTraffic	PageRank	GoogleIndex	LinksPointingToPage
\				
count	11054.000000	11054.000000	11054.000000	11054.000000
mean	0.287407	-0.483626	0.721549	0.343948
std	0.827680	0.875314	0.692395	0.569936
min	-1.000000	-1.000000	-1.000000	-1.000000
25%	0.000000	-1.000000	1.000000	0.000000
50%	1.000000	-1.000000	1.000000	0.000000
75%	1.000000	1.000000	1.000000	1.000000

max	1.000000	1.000000	1.000000	1.000000
-----	----------	----------	----------	----------

	StatsReport	class
count	11054.000000	11054.000000
mean	0.719739	0.113986
std	0.694276	0.993527
min	-1.000000	-1.000000
25%	1.000000	-1.000000
50%	1.000000	1.000000
75%	1.000000	1.000000
max	1.000000	1.000000

[8 rows x 32 columns]

df.isnull().sum()

Index	0
UsingIP	0
LongURL	0
ShortURL	0
Symbol@	0
Redirecting//	0
PrefixSuffix-	0
SubDomains	0
HTTPS	0
DomainRegLen	0
Favicon	0
NonStdPort	0
HTTPSDomainURL	0
RequestURL	0
AnchorURL	0
LinksInScriptTags	0
ServerFormHandler	0
InfoEmail	0
AbnormalURL	0
WebsiteForwarding	0
StatusBarCust	0
DisableRightClick	0
UsingPopupWindow	0
IframeRedirection	0
AgeofDomain	0
DNSRecording	0
WebsiteTraffic	0
PageRank	0
GoogleIndex	0
LinksPointingToPage	0
StatsReport	0
class	0

dtype: int64

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 11054 entries, 0 to 11053
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	Index	11054 non-null	int64
1	UsingIP	11054 non-null	int64
2	LongURL	11054 non-null	int64
3	ShortURL	11054 non-null	int64
4	Symbol@	11054 non-null	int64
5	Redirecting//	11054 non-null	int64
6	PrefixSuffix-	11054 non-null	int64
7	SubDomains	11054 non-null	int64
8	HTTPS	11054 non-null	int64
9	DomainRegLen	11054 non-null	int64
10	Favicon	11054 non-null	int64
11	NonStdPort	11054 non-null	int64
12	HTTPSDomainURL	11054 non-null	int64
13	RequestURL	11054 non-null	int64
14	AnchorURL	11054 non-null	int64
15	LinksInScriptTags	11054 non-null	int64
16	ServerFormHandler	11054 non-null	int64
17	InfoEmail	11054 non-null	int64
18	AbnormalURL	11054 non-null	int64
19	WebsiteForwarding	11054 non-null	int64
20	StatusBarCust	11054 non-null	int64
21	DisableRightClick	11054 non-null	int64
22	UsingPopupWindow	11054 non-null	int64
23	IframeRedirection	11054 non-null	int64
24	AgeofDomain	11054 non-null	int64
25	DNSRecording	11054 non-null	int64
26	WebsiteTraffic	11054 non-null	int64
27	PageRank	11054 non-null	int64
28	GoogleIndex	11054 non-null	int64
29	LinksPointingToPage	11054 non-null	int64
30	StatsReport	11054 non-null	int64
31	class	11054 non-null	int64

```
dtypes: int64(32)
```

```
memory usage: 2.7 MB
```

```
df.isnull().sum()
```

Index	0
UsingIP	0
LongURL	0
ShortURL	0
Symbol@	0
Redirecting//	0

PrefixSuffix-	0
SubDomains	0
HTTPS	0
DomainRegLen	0
Favicon	0
NonStdPort	0
HTTPSDomainURL	0
RequestURL	0
AnchorURL	0
LinksInScriptTags	0
ServerFormHandler	0
InfoEmail	0
AbnormalURL	0
WebsiteForwarding	0
StatusBarCust	0
DisableRightClick	0
UsingPopupWindow	0
IframeRedirection	0
AgeofDomain	0
DNSRecording	0
WebsiteTraffic	0
PageRank	0
GoogleIndex	0
LinksPointingToPage	0
StatsReport	0
class	0

dtype: int64

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier,
ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# Load your dataset
# Assuming your dataset is in a CSV file named 'students_data.csv'

# Assuming 'Target' is the target variable
X = df.drop('class', axis=1)
y = df['class']

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Random Forest Classifier
rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)
```

```

rf_predictions = rf_classifier.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print("Random Forest Accuracy:", rf_accuracy)

# Logistic Regression
lr_classifier = LogisticRegression(random_state=42, max_iter=1000)
lr_classifier.fit(X_train, y_train)
lr_predictions = lr_classifier.predict(X_test)
lr_accuracy = accuracy_score(y_test, lr_predictions)
print("Logistic Regression Accuracy:", lr_accuracy)

# Extra Trees Classifier
et_classifier = ExtraTreesClassifier(random_state=42)
et_classifier.fit(X_train, y_train)
et_predictions = et_classifier.predict(X_test)
et_accuracy = accuracy_score(y_test, et_predictions)
print("Extra Trees Accuracy:", et_accuracy)

# SVM Classifier
svm_classifier = SVC(random_state=42)
svm_classifier.fit(X_train, y_train)
svm_predictions = svm_classifier.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_predictions)
print("SVM Accuracy:", svm_accuracy)

Random Forest Accuracy: 0.9696969696969697
Logistic Regression Accuracy: 0.9335142469470827
Extra Trees Accuracy: 0.9687924016282226
SVM Accuracy: 0.5585707824513795

import pandas as pd
from sklearn.model_selection import StratifiedKFold
from sklearn.ensemble import RandomForestClassifier,
ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report

# Load your dataset
# Assuming your dataset is in a CSV file named 'students_data.csv'

# Define the classifiers
classifiers = {
    'Random Forest': RandomForestClassifier(),
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Extra Trees': ExtraTreesClassifier(),
    'Naive Bayes': GaussianNB()
}

```



```

# Set up k-fold stratified cross-validation
k_folds = 5 # You can adjust the number of folds
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

# Train and evaluate each classifier
for clf_name, clf in classifiers.items():
    print(f"Classifier: {clf_name}")
    accuracy_list = []
    classification_reports = []

    for train_index, test_index in skf.split(X, y):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)

        accuracy = accuracy_score(y_test, y_pred)
        accuracy_list.append(accuracy)

        classification_reports.append(classification_report(y_test,
y_pred))

# Display average accuracy and classification report
avg_accuracy = sum(accuracy_list) / k_folds
print(f"Average Accuracy: {avg_accuracy:.6f}")

print("Average Classification Report:")
for metric in classification_reports[0].split('\n')[:-1]:
    print(metric)
print("\n" + "="*50 + "\n")

```

```

Classifier: Random Forest
Average Accuracy: 0.968518
Average Classification Report:

```

	precision	recall	f1-score	support
-1	0.97	0.96	0.97	980
1	0.97	0.98	0.97	1231
accuracy			0.97	2211
macro avg	0.97	0.97	0.97	2211
weighted avg	0.97	0.97	0.97	2211

```

=====

Classifier: Logistic Regression
Average Accuracy: 0.925728
Average Classification Report:

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

-1	0.94	0.88	0.91	980
1	0.91	0.95	0.93	1231
accuracy			0.92	2211
macro avg	0.92	0.92	0.92	2211
weighted avg	0.92	0.92	0.92	2211

=====

Classifier: Extra Trees

Average Accuracy: 0.970690

Average Classification Report:

	precision	recall	f1-score	support
-1	0.97	0.96	0.97	980
1	0.97	0.98	0.97	1231
accuracy			0.97	2211
macro avg	0.97	0.97	0.97	2211
weighted avg	0.97	0.97	0.97	2211

=====

Classifier: Naive Bayes

Average Accuracy: 0.885200

Average Classification Report:

	precision	recall	f1-score	support
-1	0.84	0.93	0.88	980
1	0.94	0.86	0.89	1231
accuracy			0.89	2211
macro avg	0.89	0.89	0.89	2211
weighted avg	0.89	0.89	0.89	2211

=====

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

Specify the classifiers

```
classifiers = {
    'Random Forest': RandomForestClassifier(),
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Extra Trees': ExtraTreesClassifier(),
    'Naive Bayes': GaussianNB() # Note: SVM needs probability
                               # estimates for ROC curve
}
```

Number of folds for stratified k-fold

```

num_folds = 5

# Initialize a plot for ROC curves
plt.figure(figsize=(10, 6))

# Perform stratified k-fold validation
kf = StratifiedKFold(n_splits=num_folds, shuffle=True,
random_state=42)

for clf_name, clf in classifiers.items():
    mean_fpr = np.linspace(0, 1, 100)
    tpr_sum = 0
    auc_sum = 0

    for train_index, test_index in kf.split(X, y):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        # Fit the model
        clf.fit(X_train, y_train)

        # Get predicted probabilities for positive class
        y_probs = clf.predict_proba(X_test)[: , 1]

        # Calculate ROC curve
        fpr, tpr, _ = roc_curve(y_test, y_probs)
        tpr_sum += np.interp(mean_fpr, fpr, tpr)
        auc_sum += auc(fpr, tpr)

    # Average ROC curve over all folds
    mean_tpr = tpr_sum / num_folds
    mean_auc = auc_sum / num_folds

    # Plot the ROC curve
    plt.plot(mean_fpr, mean_tpr, label=f'{clf_name} (AUC =
{mean_auc:.2f})')

# Plot the random classifier (baseline)
plt.plot([0, 1], [0, 1], linestyle='--', color='grey', label='Random
Classifier (AUC = 0.50)')

# Customize the plot
plt.title('Receiver Operating Characteristic (ROC) Curves')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```

